

## 1 – Analysing Unoptimised Assembly Code

| Address       | Purpose                             | Value  | Size   |
|---------------|-------------------------------------|--------|--------|
| rbp           | saved base-pointer                  | rbp    | 64-bit |
| rbp - 1       | return value                        | 0 or 1 | 8-bit  |
| rbp - 8       | current                             | edi    | 32-bit |
| rbp - 12      | current                             | eax    | 32-bit |
| rbp - 16      | unused padding for 16-bit alignment | 0      | 32-bit |
| rax + rcx * 4 | “current” values in array           | edx    | 32-bit |
| r8            | array_length                        | ebx    |        |

With no optimisation, I found that the load-store addressing mode was used much more often. For each register which requires a value later, a new stack position is used to store and restore that value. For example, the addresses (rbp – 8) and (rbp – 12) both store a value of current at that point to be restored later, instead of using the same address which would be possible as eax gets the current value from edi once it is finished doing it’s calculations.

## 2 – Recursion to Iteration

| Address      | Purpose                   | Value | Size   |
|--------------|---------------------------|-------|--------|
| rbp          | saved base-pointer        | rbp   | 64-bit |
| rbp - 8      | saved rbx                 | rbx   | 64-bit |
| r9 + rbx * 4 | “current” values in array | edi   | 32-bit |
| r8           | array_length              | ebx   | 32-bit |

For all optimisation levels, more stack space was required than the red-zone, as space were required to store the array and it’s values in the stack.

With the higher optimisation, the biggest difference in the stack creation is the usage of load-store addressing modes was limited as much as possible. This is because register-register addressing is much faster. This can be seen as the temporary current values are stored in edi and edx where one stores the current value whilst the other does calculations on it.

The function `collatz_recurse` function does not appear with this optimisation because the current value for `collatz_recurse` is set using the start value in `collatz`, and so the compiler just uses the initial argument and changes that using the two loops of `current * in_mult + in_add` and `current / in_div` without needing to set or pass through variables to a second function.

I also noticed that there are certain commands which are skipped because they have no affect on the code. On line 136 and 140 of `collatz0-x86-commented.s`, there is a command “`andb $1, %al`” which does an “and” logical operation on `%al`, which means that `%al` is 1 when `1 & %al` are 1, and 0 with any

other values. As 1 is always 1,  $\%al = 1$  when  $\%al = 1$  which means that  $\%al = \%al$ . This is not included in the optimised version.

Also, on line 141, the command “movzbl  $\%al$ ,  $\%eax$ ” is skipped in the optimised version because the register  $\%al$  is just the first 8-bits of the 32-bit register  $\%eax$ . Setting the 32-bit register to the 8-bit register will just mean that  $\%eax = \%eax$ .

### 3 – ARMv8

| Command                            | Meaning                                                                                                       | Addressing Mode |
|------------------------------------|---------------------------------------------------------------------------------------------------------------|-----------------|
| stp x29, x30, [sp, #-16]!          | Stores values x29 and x30 into the stack space in space between stack pointer (sp) and $sp - 16$ .            | Relative        |
| mov x29, sp                        | Stores the second value sp into the address x29                                                               | Register        |
| adrp x8, array_length              | Gets the label address (not the value) of array_length and stores it in x8                                    | Relative        |
| ldrsx x9, [x8, :l012:array_length] | Loads the signed word value of array_length by calculating the address and then gets the value, storing in x9 | Register        |
| cmp w9, #199                       | Subtracts 199 from w9 and stores flags without affecting w9                                                   | Immediate       |
| b.le .LBB1_2                       | Checks the flags for signed less-than or equal to 0 and branches (jumps) to .LBB1_2                           | Relative        |
| add w11, w9, #1                    | Adds w9 and the number 1 to w11, storing in w11                                                               | Immediate       |
| str w0, [x10, x9, lsl #2]<br>sdiv  | Stores the value w0 in the stack at $x10 + x9 * 2$                                                            | Immediate       |

In ARM, the destination and source is often swapped compared to x86, where the destination comes first.

The calling convention can be seen in the storing of x29 and x30 from the stack which correspond to the frame-pointer base-pointer and link-register which are stored at the beginning of the function.

The different addressing modes used were register, which uses the address values stored in a register (such as the ldrsw above); immediate, which use fixed number values (such as the cmp above) without needing to look in any registers; and relative (such as stp above) where the address is calculated using a mix of register value, as well as an offset from that register address.