

# CS3101 - Databases Exam

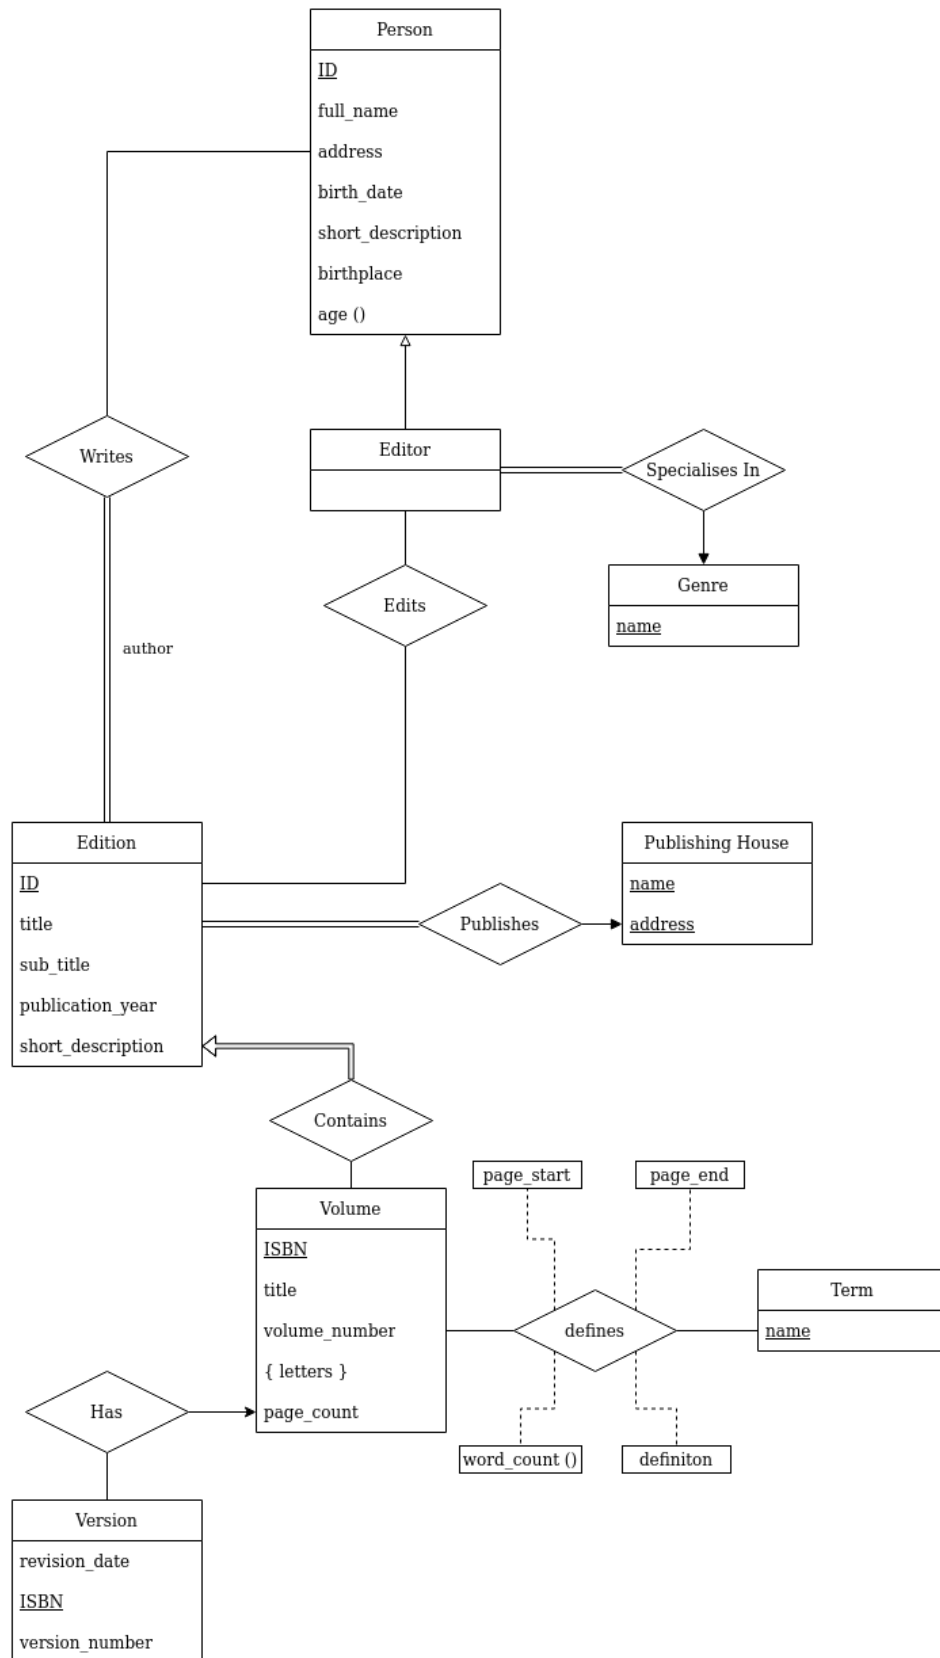
200007413

22 Apr 2023

## **1 TODO**

### **1.a**

For this scenario, I assumed that the ISBN are unique across all volumes, editions, and versions.



## 1.b

person(id, full\_name, address, birth\_date, short\_description, birthplace)

editor\_specialty(person\_id\*, genre\_name)

writes(person\_id\*, edition\_id\*)

edits(person\_id\*, edition\_id\*)

edition(id, title, sub\_title, publication\_year, short\_description)

publishing\_house(name, address)

edition\_publisher(publisher\_name, address)

## 1.c

### 1.c.1

```
SELECT id, publication_year FROM edition
```

### 1.c.2

```
SELECT e.id, COUNT(v.ISBN) AS count FROM edition as e LEFT JOIN volume as v ON  
v.edition_id = e.id WHERE e.publication_year < 1850 GROUP BY e.edition_id
```

### 1.c.3

```
SELECT start_page, end_page FROM term_definition AS td LEFT JOIN volume AS v ON  
td.volume_ISBN = v.ISBN WHERE td.name = "abacus" AND exists (SELECT * FROM edition  
as e WHERE e.id = v.edition_id AND e.id = 1)
```

## 2

### 2.a

#### 2.a.1

```
student = { id → name
            }
tutors = { tutor_id → tutor_email
           }
practical_grade = { student_id,module_id,practical_number → grade
                   }
```

#### 2.a.2

At the given exact moment, there would be a functional dependency between `module_id` and `practical_number`, as each value of `module_id` refers to a unique `practical_number`.

However, this functional dependency will not exist once a second practical for a module is added, as then one module is not guaranteed to reference a unique `practical_id`.

### 2.b

#### 2.b.1

Yes, the *tutors* relation is in first normal form, as all its attributes are atomic. Each tuple refers to a single student using `student_id`, and refers to a single staff member using `tutor_id`, with a single email.

#### 2.b.2

Yes, because *tutors* is in first normal form, and each of the attributes are either candidate keys or full functionally dependent on a candidate key.

Both `tutor_id` and `student_id` are required to uniquely identify the relationship, and so are candidate keys, and then `tutor_email` is functionally dependent on the `tutor_id`.

#### 2.b.3

Under the assumption that tutors do not have more than one email address, *tutors* is not in third normal form.

Whilst it is in second normal form, and `tutor_id` and `student_id` are part of the super-key, `tutor_id` itself is not a superkey, and there is a functional dependency from `tutor_id` to the non-candidate-

key attribute `tutor_email`. Additionally, `tutor_email-tutor_id = tutor_email`, which is itself not in a candidate key.

#### 2.b.4

This is not in Boyce-Codd normal form (BCNF) as it is not in third-normal form.

The condition for BCNF is the same as 3NF in that every functional dependency  $\alpha \rightarrow \beta$  must either be trivial or be dependent on a superkey, with the added restriction that  $\beta - \alpha$  being in a candidate key is not sufficient. As the `tutor_id → tutor_email` already doesn't meet the less strict 3NF conditions, it does not meet BCNF.

#### 2.c

Artificial primary keys are those which do not relate to any real data, but are added to ensure that there is a unique identifier for each entry, instead of relying upon the three natural candidate keys to uniquely identify each tuple.

The advantages to using the artificial ID value is that you can create queries using just that identifier, instead of requiring that someone know the `module_id`, `student_id`, and `practical_number`. Practically, this could allow students to access their practical grades just with the `practical_grade` ID.

The downside to using the ID attribute is it will add redundant information, and prevent the schema from being in Boyce-Codd normal form. This is because each entry is already uniquely identified using the three candidate keys, and so the ID is functionally dependent on these and visa-versa.

This redundancy would add additional space requirements for storing the table, as a new not-strictly required column is added to each tuple.

Another risk from artificial keys is with the chance for duplicate entries. If ID is set as the sole unique superkey, then more than one entry could point to the same "useful" data, but just with a new ID, practically adding duplicate data. Without artificial keys, only unique combinations of `student_id`, `module_id`, and `practical_number` can be added.

#### 2.d TODO

### 3

#### 3.a

A *mediumblob/longblob* type may be appropriate for *user.picture* for storing the raw picture data itself. This would be a large data type, and would store the picture in the database.

Another appropriate type could be a *varchar*, which would contain the URL to that user's picture.

The benefit of a blob data type is that no additional calls need to be made to retrieve the users profile picture, as the call itself returns it's raw data.

However, this means the database will be much larger than having a varchar containing a url.

Another advantage of the varchar URL method is the image can be stored at the URL as an image file. Storing as a blob ignores the file tyoe, meaning the program accessing the data must interpret the image properly. Storing the image as a file can prevent the need to interpret raw data in this way.

### 3.b

#### 3.b.1

$$\Pi_{\text{username}}(\sigma_{\text{rating} > 1800}(\text{user}))$$

#### 3.b.2

$$\Pi_{\text{p1\_username}}(\text{game} \bowtie \sigma_{\text{name} = \text{"Sammy Jacobs Cup"}}(\text{tournament}))$$

### 3.c

SELECT UNIQUE p1\_username FROM game WHERE EXISTS (SELECT tournament\_id FROM tournament WHERE tournament\_id = id AND tournament.name = "Sammy Jacobs Cup")

### 3.d

#### 3.d.1

manutdlover

joker77,manudlover,tobym

alan34,alexr,bigbob joker77,kingkhalid,md216 mudlover,ohio\_alice,pkg113 tobym

### **3.d.2**

Increasing the value of N would allow for more items to be added before the tree needs to be reorganized.

An issue with this is that the graph would become more uneven, thus reducing the efficiency when searching.

### **3.d.3**

The advantage of a hash index is that exact match queries are much faster. An example of a query which would run faster for a hash index would be:

```
SELECT * FROM user WHERE username = "bob"
```

Queries which have non-exact matching criteria would be faster for B+ indexes. For example:

```
SELECT * FROM user WHERE username LIKE "bo%son"
```